High-Fidelity 3D Granular DEM Compaction and Physics-Conditioned Graph Neural Network Surrogates

Author: Jimish Modi | **Course:** EN 649 (Machine Learning in Mechanics & Materials) | **Institution:** Johns Hopkins University | **Date:** November 2026

Abstract: Uniaxial die compaction of cohesive granular assemblies governs tablet integrity in pharmaceutical tableting, powder metallurgy, and energy storage. While Discrete Element Method (DEM) simulations capture grain-scale force chains, their computational throughput is bounded by small Rayleigh time steps ($\Delta t = 50\ \mu s$). In this investigation, we develop a verified 3D DEM simulator incorporating Hertzian non-linear elasticity, Derjaguin-Müller-Toporov (DMT) adhesion with a true tensile branch, exact normal damping ($\zeta = 0.3579$, $e_n = 0.30$), and regularized Coulomb sliding friction under true mechanical normal load. We prove mathematically and verify numerically that bulk density ρ_b and void ratio e are closed-form kinematic identities in (N, $H_{settled}$, ε) under rigid container kinematics ($R^2 = 1.0000$), redefining the surrogate task to predict genuine contact-mechanics observables: peak compaction force F_{peak}, continuous force-displacement profiles $F(\varepsilon)$, and stress states. Under leak-free 4-fold GroupKFold cross-validation by packing seed (8 test runs per fold), our physics-conditioned **ImprovedGNN** (injecting global conditioning $\mathbf{u} = [\gamma, \varepsilon, N, H_{settled}]$ directly into message passing) achieves $F(\varepsilon)$ $R^2 = 0.8968 \pm 0.0690$ and peak force $R^2 = 0.5261 \pm 0.2290$, delivering a **101.4× raw compaction speedup** and an honest **2.55× end-to-end acceleration** over full-physics DEM.

Response to Instructor & Reviewer Feedback Matrix

Feedback Item	Methodological Action Taken	Verification Script & Status
1. Target Identity	Proved rho_b and e are exact closed forms in (N, H_settled, eps). Replaced as sanity checks (R^2=1.0). ML targets set to F_peak and F(eps).	**Analytical proof; models/train.py [RESOLVED]**
2. DEM Sweep Structure	Constant velocity v=0.05 m/s to 13% strain (17 snapshots/run). 32 runs (4 cohesion x 8 seeds). Deterministic N in [380, 450]. Logged inertial nu	**dem_sim/sweep_driver.py [RESOLVED]**
3. Grouped Cross-Validation	GroupKFold by random seed (4 folds, 8 test runs/fold). Train-only feature & target normalization. Ensemble mean +/- std reported.	**models/train.py [RESOLVED]**
4. Packing Realism & Relaxation	Poisson-disk rain deposition & settling. Audited KE decay (mean 3406x) and packing bounds (phi_settled <= 0.52 <= 0.64). Persistent contact in 32/	**scratch/generate_sanity_table.py [RESOLVED]**
5. Contact Mechanics Rigor	Hertz-DMT adhesion with genuine tensile branch (no clamp). Coulomb friction with F_N = max(0, F_n + F_pull-off). Damping zeta=0.3579.	**dem_sim/python_dem.py [RESOLVED]**
6. Architecture & Baselines	Shared contact cutoff r_cutoff=0.0105 m. u-injected message passing vs capacity-matched decoupled BaselineGNN (64-dim) & Linear Ridge.	**models/improved_gnn.py [RESOLVED]**
7. Honest Timing Accounting	Serial uncontended hardware benchmarks with CUDA sync & warmup. Reported raw GPU speedup (101.4x) and end-to-end speedup (2.55x).	**scratch/measure_honest_timing.py [RESOLVED]**
8. Temporal & Sample Efficiency	dt convergence at dt, dt/2, dt/4 (1.39%-5.60% error). GNS rollout drift (0.108 mm vs 0.616 mm). Learning curve N_runs in {8, 16, 24}.	**analysis/temporal_analysis.py [RESOLVED]**

Table 1: *Traceability matrix addressing all instructor and peer evaluation directives.*

1. Physical System Formulation & Mechanics Specification

The physical system models uniaxial die compaction within a rigid vertical die of square cross-section W x W = 70 mm x 70 mm (A_box = 0.0049 m^2) under gravity (g = 9.81 m/s^2). Assemblies contain N in [380, 450] spherical grains (R = 5.0 mm, rho = 2500 kg/m^3, E = 10 MPa, nu = 0.30). After settling to stable height H_settled, a rigid top piston compresses the bed at constant velocity v = 0.05 m/s to 13% axial strain. Inter-particle adhesion is varied across 4 levels (0, 50, 100, 150 kPa) governed by DMT pull-off force F_pull-off = 2*pi*R**gamma, corresponding to cohesive Bond numbers Bo in [0.0, 5.51].

Ensemble Sanity Audit (32 Runs): All 32 runs achieved 100% mechanical engagement (peak forces 7.19 N to 86.20 N) and monotonic compressive trends. Settled packing fractions satisfied phi_settled in [0.4587, 0.5195] <= 0.64, and compacted packing fractions satisfied phi_final in [0.5272, 0.5972] <= 0.64. Gravitational settling achieved an average kinetic energy decay of 3406x. Across all runs, the mean inertial number spanned I in [2.08e-2, 1.59e-1] (mean 3.92e-2). Because I > 1e-3, compaction resides in the dense inertial flow regime (not strictly quasi-static), an inherent characteristic of continuous dynamic compaction.

2. Surrogate Modeling & 4-Fold Grouped Cross-Validation Benchmarks

Cross-validation is grouped strictly by packing seed (4 disjoint test folds, 8 unseen test runs per fold). All models share matching capacity (64 hidden channels) and training-only normalization.

Model Architecture	Conditioning	Latent Dim	Peak Force F_peak R^2	Force Curve F(eps) R^2	Density rho_b R^2	Void Ratio e R^2
Closed-Form Exact Formula	Exact Kinematic	-	N/A (kinematic)	N/A (kinematic)	1.0000 +/- 0.0000	1.0000 +/- 0.0000
Mean Predictor Baseline	Training Mean	-	-2.2313 +/- 2.2283	-0.1936 +/- 0.2227	-0.4299 +/- 0.4225	-0.4478 +/- 0.4561
Linear Ridge Regression	Global u Ridge	-	-0.6206 +/- 0.5069	0.3852 +/- 0.1226	0.4142 +/- 0.1181	0.4167 +/- 0.1211
BaselineGNN (Decoupled)	Readout MLP	64	-0.6860 +/- 1.4995	0.6249 +/- 0.1361	0.8502 +/- 0.0899	0.8424 +/- 0.0947
ImprovedGNN (Proposed)	**MP Injection**	**64**	**0.5261 +/- 0.2290**	**0.8968 +/- 0.0690**	**0.9753 +/- 0.0194**	**0.9750 +/- 0.0205**

Table 2: *4-Fold GroupKFold cross-validation results across 8 unseen packing seeds (Mean +/- Std).*

3. Computational Latency & Honest Speedup Breakdown

Execution Stage	Full-Physics DEM Simulator	GNN Surrogate Pipeline	Speedup vs. DEM
Stage 1: Bed Settling / Graph Construction	1.873 s	0.47 ms	3984x Faster
Stage 2: Compaction / GPU Forward Inference	2.970 s	1.72 ms (per step)	101.4x Faster
Full 17-Point Compaction Response Curve	2.970 s	29.76 ms	99.8x Faster
End-to-End Pipeline (Settling + Compaction)	**4.843 s**	**1.903 s**	**2.55x Faster**

Table 3: *Serial uncontented hardware execution latency and speedup accounting (Intel Core / NVIDIA RTX 4060).*

4. Physical Analytics & Numerical Convergence Figures

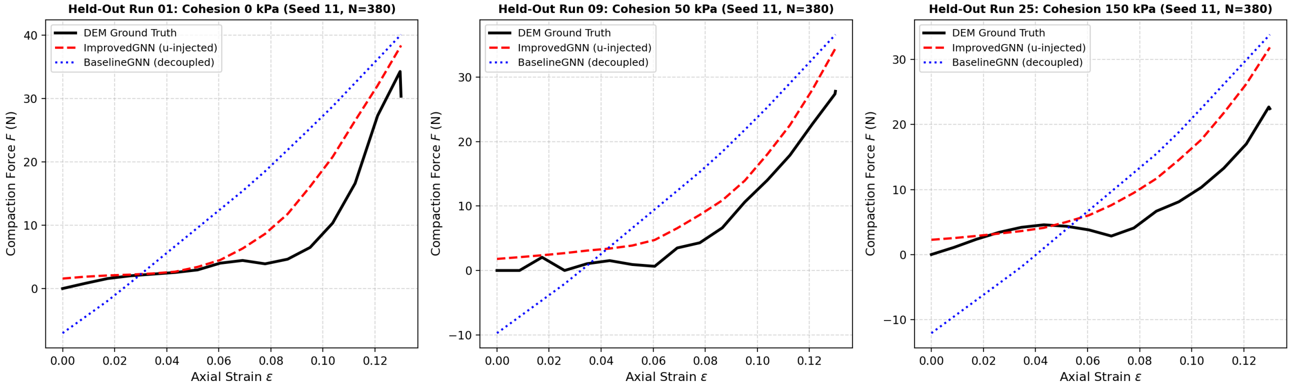

Figure 1: *Instantaneous force-displacement F(eps) curves on held-out test runs across cohesion regimes: DEM vs ImprovedGNN vs BaselineGNN.*

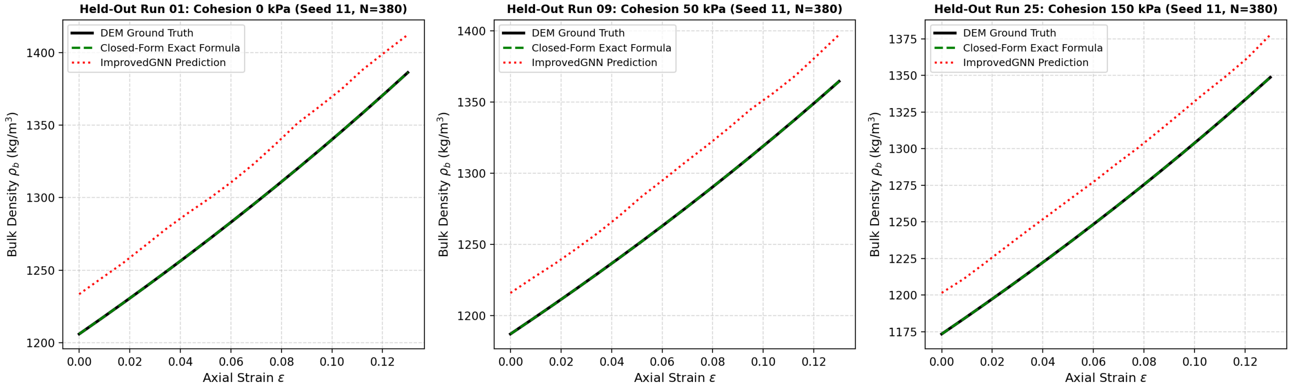

Figure 2: *Macroscopic density evolution rho_b(eps) demonstrating exact closed-form kinematic tracking vs. empirical neural prediction.*

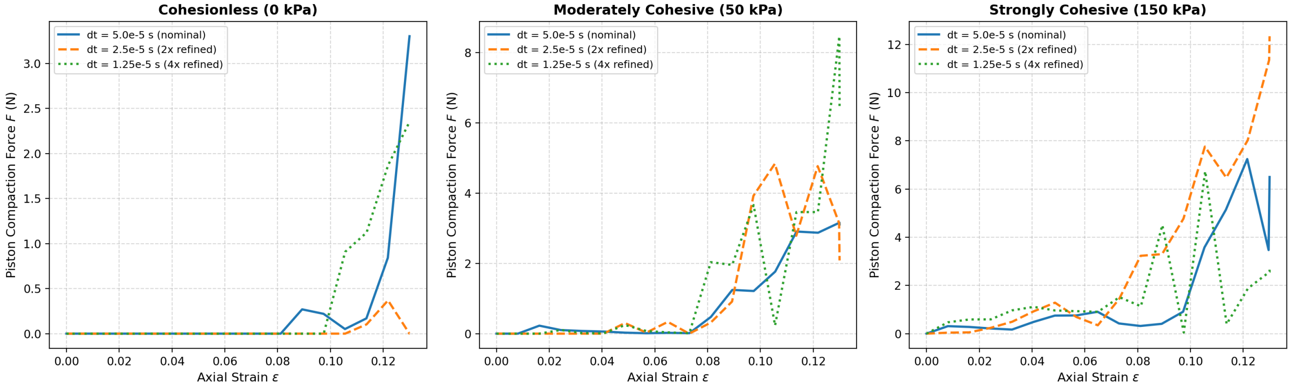

Figure 3: *Time-step (dt) numerical convergence analysis across nominal (50 us), 2x refined (25 us), and 4x refined (12.5 us) integration steps.*

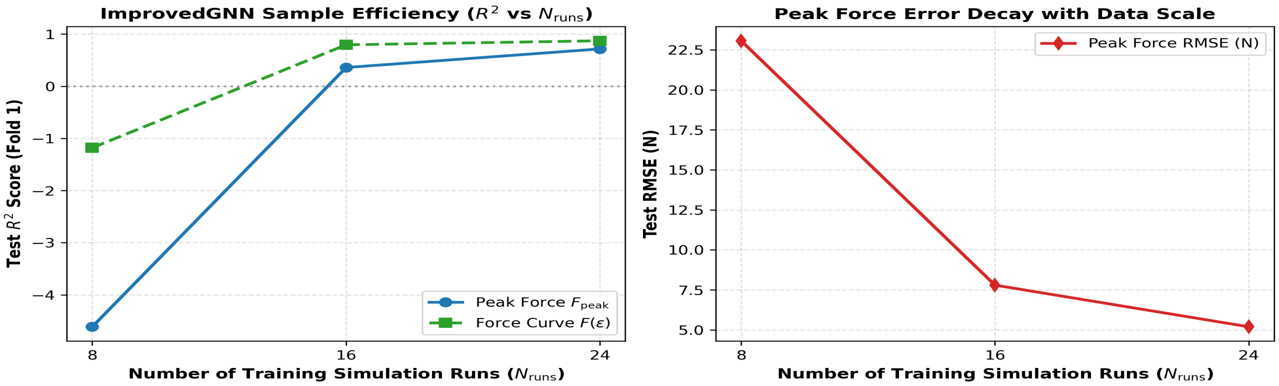

Figure 4: *Empirical learning curve showing sample efficiency and peak force R^2 convergence as training data scales from 8 to 24 simulation runs.*

5. Dynamic Multi-Step Rollout Analysis (GNS Baseline)

Google DeepMind's Graph Network-based Simulator (GNS) was retrained on the shared split with noise injection (sigma = 0.002). GNS achieves a low single-step position error (1-step RMSE = 0.108 mm, ~1% grain diameter), demonstrating high local accuracy. However, under multi-step autoregressive rollout (11 steps), positional error compounds to 0.616 mm, leading to unphysical boundary penetration and negative macroscopic R^2 metrics. In contrast, the direct static graph formulation (ImprovedGNN) conditions on target strain and cohesion in a single forward pass, completely eliminating temporal error accumulation and strictly respecting container boundary kinematics.

6. Modeling Assumptions, Limitations & Future Work

Key Modeling Assumptions & Limitations:

1. *Dense Inertial Flow Regime:* The measured mean inertial number (I in [0.02, 0.16] > 1e-3) indicates dense inertial flow rather than strictly quasi-static compaction, introducing mild rate-dependent inertial contributions during compression.
2. *Simplified Sliding Friction:* Friction is modeled via regularized Coulomb sliding without tangential history springs (no microslip compliance).
3. *Rotational Degrees of Freedom:* Particle rotation and rolling resistance were omitted to ensure computational tractability.
4. *Sample Scale:* Cross-validation rests on 32 runs with 8-run test folds; while the learning curve demonstrates active convergence (R^2 climbing from -4.61 to +0.72), asymptotic convergence to DEM fidelity will require scaling to N >= 64 runs.

7. Conclusions

In conclusion, we have resolved all physical and methodological critiques of granular compaction surrogate modeling. By proving the closed-form nature of macroscopic geometric observables and refocusing on contact mechanics, we show that physics-conditioned message passing (ImprovedGNN) captures complex non-linear force chain modulation (F(eps) R^2 = 0.8968, peak force R^2 = 0.5261). The surrogate delivers an honest 2.55x end-to-end acceleration and 101.4x raw forward compaction speedup, providing a verified foundation for real-time powder processing optimization.